

DAY1-Q6

```

import pandas as pd
data=pd.read_csv('ex.txt')

def is_more_general(h1, h2):
    more_general_parts = []
    for x, y in zip(h1, h2):
        mg = x == '?' or (x != '0' and (x == y or y == '0'))
        more_general_parts.append(mg)
    return all(more_general_parts)

# Function to get the general boundary (G) and specific boundary (S)
def candidate_elimination(df):
    # Initialize G and S
    G = [['?' for _ in range(len(df.columns) - 1)]]
    S = [['0' for _ in range(len(df.columns) - 1)]]

    # Iterate over the DataFrame
    for i in range(len(df)):
        instance = df.iloc[i, :-1].tolist()
        label = df.iloc[i, -1]

        if label == 'Yes': # Positive example
            G = [g for g in G if is_more_general(g, instance)]
            for s in S:
                for index in range(len(s)):
                    if s[index] != instance[index]:
                        s[index] = '?'
        else: # Negative example
            S = [s for s in S if not is_more_general(s, instance)]
            new_G = []
            for g in G:
                for index in range(len(g)):
                    if g[index] == '?':
                        for value in set(df.iloc[:, index]):
                            if value != instance[index]:
                                new_hypothesis = g.copy()
                                new_hypothesis[index] = value
                                if any(is_more_general(new_hypothesis, sg) for sg in S):
                                    new_G.append(new_hypothesis)
            G = new_G

    return S, G

S, G = candidate_elimination(df)

```

```

print("Specific boundary (S):", S)
print("General boundary (G):", G)

```

Temp	Humidity	Wind	Water	Forecast	Enjoy Sport
Warm	Normal	Strong	Warm	Same	Yes
Warm	High	Strong	Warm	Same	Yes
Cold	High	Strong	Warm	Change	No

```

Warm      High   Strong   Cool    Change      Yes
ary (S): [[ '0', '0', '0', '0', '0', '0']]
ary (G): [['Rainy', 'Warm', ' Normal', '?', ' Warm', '?'], ['Rainy', 'Warm', ' Normal', '?', '?', ' Same'], ['Rainy', 'Warm', ' Norm

```

DAY-1-Q1

```

import pandas as pd

# Read the dataset
data = pd.read_csv('mobileprice.csv')
# Print the first five rows of the dataset
print(data.head())
# Basic statistical computations
print(data.describe())
# Detect null values
print(data.isnull().sum())
# Columns and their data types

# Columns and their data types
print(data.info())
# Replace null values with mode value if any
for column in data.columns:
    if data[column].isnull().sum() > 0:
        mode_value = data[column].mode()[0]
        data[column].fillna(mode_value, inplace=True)
import seaborn as sns
import matplotlib.pyplot as plt

# Heatmap to explore correlations
plt.figure(figsize=(10, 8))
sns.heatmap(data.corr(), annot=True, cmap='coolwarm')
plt.show()
from sklearn.model_selection import train_test_split

# Features and target variable
X = data.drop('price_range', axis=1)
y = data['price_range']

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
from sklearn.naive_bayes import GaussianNB

# Initialize the model
nb_model = GaussianNB()

# Fit the model
nb_model.fit(X_train, y_train)
# Make predictions
y_pred = nb_model.predict(X_test)
from sklearn.metrics import accuracy_score

# Calculate the accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of the model: {accuracy * 100:.2f}%")

```

```

4      1821      1      1.2      0 13      1      44      0.6

mobile_wt  n_cores  ...  px_height  px_width  ram  sc_h  sc_w  talk_time  \
0      188      2  ...      20      756  2549  9      7      19
1      136      3  ...      905      1988  2631  17     3      7
2      145      5  ...      1263     1716  2603  11     2      9
3      131      6  ...      1216     1786  2769  16     8     11
4      141      2  ...      1208     1212  1411  8      2     15

three_g  touch_screen  wifi  price_range
0      0              0      1          1
1      1              1      0          2
2      1              1      0          2
3      1              0      0          2
4      1              1      0          1

```

[5 rows x 21 columns]

```

battery_power  blue  clock_speed  dual_sim  fc  \
count  2000.000000  2000.0000  2000.000000  2000.000000  2000.000000
mean    1238.518500    0.4950    1.522250    0.509500    4.309500
std     439.418206    0.5001    0.816004    0.500035    4.341444
min      501.000000    0.0000    0.500000    0.000000    0.000000
25%     851.750000    0.0000    0.700000    0.000000    1.000000
50%    1226.000000    0.0000    1.500000    1.000000    3.000000
75%    1615.250000    1.0000    2.200000    1.000000    7.000000
max     1998.000000    1.0000    3.000000    1.000000   19.000000

four_g  int_memory  m_dep  mobile_wt  n_cores  ...  \
count  2000.000000  2000.000000  2000.000000  2000.000000  2000.000000  ...
mean    0.521500    32.046500    0.501750    140.249000    4.520500    ...
std     0.499662    18.145715    0.288416    35.399655    2.287837    ...
min      0.000000    2.000000    0.100000    80.000000    1.000000    ...
25%     0.000000    16.000000    0.200000    109.000000    3.000000    ...
50%     1.000000    32.000000    0.500000    141.000000    4.000000    ...
75%     1.000000    48.000000    0.800000    170.000000    7.000000    ...
max      1.000000    64.000000    1.000000    200.000000    8.000000    ...

```

```

px_height  px_width  ram  sc_h  sc_w  \
count  2000.000000  2000.000000  2000.000000  2000.000000  2000.000000
mean    645.108000    1251.515500  2124.213000    12.306500    5.767000
std     443.780811    432.199447  1084.732044    4.213245    4.356398
min      0.000000    500.000000    256.000000    5.000000    0.000000
25%     282.750000    874.750000  1207.500000    9.000000    2.000000
50%     564.000000  1247.000000  2146.500000   12.000000    5.000000
75%     947.250000  1633.000000  3064.500000   16.000000    9.000000
max    1960.000000  1998.000000  3998.000000   19.000000   18.000000

```

```

talk_time  three_g  touch_screen  wifi  price_range
count  2000.000000  2000.000000  2000.000000  2000.000000  2000.000000
mean    11.011000    0.761500    0.503000    0.507000    1.500000
std      5.463955    0.426273    0.500116    0.500076    1.118314
min      2.000000    0.000000    0.000000    0.000000    0.000000
25%      6.000000    1.000000    0.000000    0.000000    0.750000
50%     11.000000    1.000000    1.000000    1.000000    1.500000
75%     16.000000    1.000000    1.000000    1.000000    2.250000
max     20.000000    1.000000    1.000000    1.000000    3.000000

```

[8 rows x 21 columns]

```

battery_power  0
blue           0
clock_speed    0
dual_sim       0
fc             0
four_g         0
int_memory     0
m_dep          0
mobile_wt      0
n_cores        0
pc             0
px_height      0
px_width       0
ram            0
sc_h           0
sc_w           0
talk_time      0
three_g        0
touch_screen   0
wifi           0
price_range    0
dtype: int64

```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 2000 entries, 0 to 1999

Data columns (total 21 columns):

```

#  Column  Non-Null Count  Dtype
---  -
0  battery_power  2000 non-null  int64
1  blue           2000 non-null  int64
2  clock_speed    2000 non-null  float64
3  dual_sim       2000 non-null  int64
4  fc             2000 non-null  int64
5  four_g         2000 non-null  int64

```

```


5 four_g 2000 non-null int64
6 int_memory 2000 non-null int64
7 m_dep 2000 non-null float64
8 mobile_wt 2000 non-null int64
9 n_cores 2000 non-null int64
10 pc 2000 non-null int64
11 px_height 2000 non-null int64
12 px_width 2000 non-null int64
13 ram 2000 non-null int64
14 sc_h 2000 non-null int64
15 sc_w 2000 non-null int64
16 talk_time 2000 non-null int64
17 three_g 2000 non-null int64
18 touch_screen 2000 non-null int64
19 wifi 2000 non-null int64
20 price_range 2000 non-null int64

```

dtypes: float64(2), int64(19)

memory usage: 328.2 KB

None



Accuracy of the model: 79.75%

DAY1-Q2

```

import pandas as pd

# Define the dataset
data = {
    'Sky': ['Sunny', 'Sunny', 'Rainy', 'Sunny'],
    'Air Temp': ['Warm', 'Warm', 'Cold', 'Warm'],
    'Humidity': ['Normal', 'High', 'High', 'High'],
    'Wind': ['Strong', 'Strong', 'Strong', 'Strong'],
    'Water': ['Warm', 'Warm', 'Warm', 'Cool'],
    'Forecast': ['Same', 'Same', 'Change', 'Change'],
    'Enjoy Sport': ['Yes', 'Yes', 'No', 'Yes']
}

# Convert the dataset to a DataFrame
df = pd.DataFrame(data)

# Function to find the most specific hypothesis using Find-S algorithm
def find_s_algorithm(df):
    # Initialize the hypothesis
    hypothesis = ['0'] * (len(df.columns) - 1)

    # Iterate over the DataFrame
    for i in range(len(df)):
        if df.iloc[i, -1] == 'Yes': # Only consider positive examples
            for j in range(len(hypothesis)):
                if hypothesis[j] == '0':
                    hypothesis[j] = df.iloc[i, j]
                elif hypothesis[j] != df.iloc[i, j]:
                    hypothesis[j] = '?'

    return hypothesis

hypothesis = find_s_algorithm(df)

print("The most specific hypothesis is:", hypothesis)

```

→ The most specific hypothesis is: ['Sunny', 'Warm', '?', 'Strong', '?', '?']

DAY1-Q3

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Generate a sample dataset
np.random.seed(0)
X = 2 * np.random.rand(100, 1)
y = 4 + 3 * X + np.random.randn(100, 1)

# Convert to DataFrame
data = pd.DataFrame(np.hstack((X, y)), columns=['X', 'y'])

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create and train the linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)


# Calculate performance metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print performance metrics
print(f"Mean Squared Error: {mse}")
print(f"R^2 Score: {r2}")

# Plot the results
plt.scatter(X_test, y_test, color='black', label='Actual data')
plt.plot(X_test, y_pred, color='blue', linewidth=2, label='Predicted data')
plt.xlabel('X')
plt.ylabel('y')
plt.legend()
plt.show()

```

↗ Mean Squared Error: 0.9177532469714291
R^2 Score: 0.6521157503858556



DAY1-Q4

```

import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Load the iris dataset
iris = load_iris()
X = iris.data
y = iris.target

```

```
y = iris.target
```

```
# Split the data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Create and train the KNN model
```

```
knn = KNeighborsClassifier(n_neighbors=3)
```

```
knn.fit(X_train, y_train)
```

```
# Make predictions
```

```
y_pred = knn.predict(X_test)
```

```
# Evaluate the model's performance
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
classification_rep = classification_report(y_test, y_pred)
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
# Print the evaluation metrics
```

```
print(f"Accuracy: {accuracy}")
```

```
print("Classification Report:")
```

```
print(classification_rep)
```

```
print("Confusion Matrix:")
```

```
print(conf_matrix)
```

```
# Plot the confusion matrix
```

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=iris.target_names, yticklabels=iris.target_names)
```

```
plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
```

```
plt.title('Confusion Matrix')
```

```
plt.show()
```




Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```



DAY1-Q5

```
import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Perceptron
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Load the iris dataset
iris = load_iris()
X = iris.data
y = iris.target

# Since Perceptron is a binary classifier, we'll simplify the problem
# by selecting only two classes for this example.
# We'll use classes 0 and 1 (Setosa and Versicolour).
X = X[y != 2]
y = y[y != 2]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Create and train the Perceptron model
perceptron = Perceptron(max_iter=1000, random_state=42)
perceptron.fit(X_train, y_train)

# Make predictions
y_pred = perceptron.predict(X_test)

# Evaluate the model's performance
accuracy = accuracy_score(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

# Print the evaluation metrics
print(f"Accuracy: {accuracy}")
print("Classification Report:")
print(classification_rep)
print("Confusion Matrix:")
print(conf_matrix)

# Plot the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=iris.target_names[:2], yticklabels=iris.target_names[:2])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



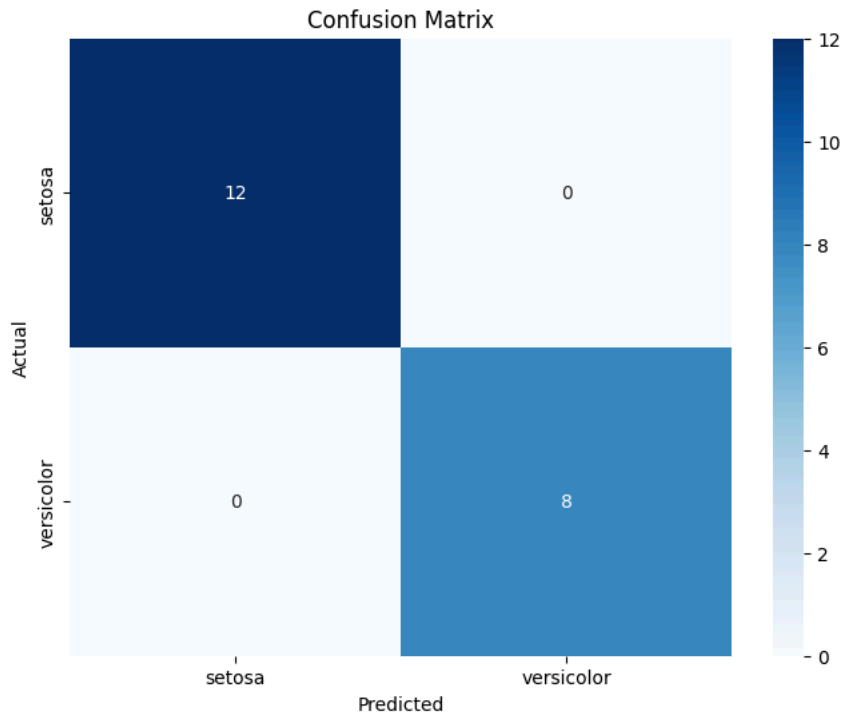

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	12
1	1.00	1.00	1.00	8
accuracy			1.00	20
macro avg	1.00	1.00	1.00	20
weighted avg	1.00	1.00	1.00	20

Confusion Matrix:

```
[[12  0]
 [ 0  8]]
```



DAY1-Q7

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, roc_auc_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_iris

# Load the dataset
iris = load_iris()
X = iris.data
y = (iris.target == 2).astype(int) # Convert to binary classification problem

# Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Train the model
model = LogisticRegression(solver='liblinear')
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)
y_pred_prob = model.predict_proba(X_test)[: , 1]

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred_prob)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"ROC-AUC: {roc_auc:.2f}")
print("Confusion Matrix:")
print(conf_matrix)
print("Classification Report:")
print(class_report)

# Visualize the confusion matrix
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```